

# CS 278: Computational Complexity Theory

## Homework 2

Due: **October 24th, 2025**

Fall 2025

### Instructions:

- Collaboration is allowed but solutions must be written independently. List collaborators and external resources.
- Write your solutions in L<sup>A</sup>T<sub>E</sub>X and submit a single PDF by email to `ljiexchen@berkeley.edu` with subject “CS 278 - Homework 2 – [Your Name]”.
- Name your file `CS278-HW2-[YourName].pdf`. **Deadline:** 11:59pm Pacific Time October 24th, 2025.
- Late submissions lose **10%** per day (e.g., three days late →  $0.9^3$  of your score).
- This homework has 4 problems, each with multiple parts, totaling **160 points**. As in HW1, our course policy on homework aggregation applies unchanged.

# 1 Problem 1 (40 pts): A circuit lower bound for quadratic space

Show that languages decidable in quadratic space do not admit linear-size circuits.

**Statement.** Prove that  $\text{SPACE}[n^2] \not\subseteq \text{SIZE}(O(n))$ . In other words, there is a language  $L \in \text{SPACE}[n^2]$  such that every Boolean circuit family computing  $L$  on inputs of length  $n$  has size  $\omega(n)$  for all sufficiently large  $n$ .

**2 Problem 2 (40 pts):  $\text{NEXP} \subseteq \text{coNEXP}_{/n+1}$** 

**Statement.** Prove that  $\text{NEXP} \subseteq \text{coNEXP}_{/n+1}$ . In other words, for every language  $L \in \text{NEXP}$ , prove that there exists an advice function  $a(n) \in \{0, 1\}^{n+1}$  such that there exists a coNEXP machine  $M$  that, for every  $n \in \mathbb{N}$ , given the correct advice  $a(n)$ ,  $M$  computes  $L$  on all  $n$ -bit inputs.

### 3 Problem 3 (40 pts): Prove $\text{CL} \subseteq \text{ZPP}$

The class CL (catalytic logspace) allows an algorithm logarithmic *clean* space and polynomially many *dirty* bits that must be restored at the end. Show that every  $L \in \text{CL}$  has a zero-error expected-polynomial-time algorithm.

**Definition of ZPP.** The complexity class ZPP (Zero-error Probabilistic Polynomial time) consists of all languages  $L$  for which there exists a probabilistic polynomial-time Turing machine  $M$  such that:

1. For every input  $x$ ,  $M(x)$  outputs either 0, 1, or  $\perp$  (“don’t know”).
2. If  $x \in L$ , then  $\Pr[M(x) = 1] \geq \frac{1}{2}$  and  $\Pr[M(x) = 0] = 0$ .
3. If  $x \notin L$ , then  $\Pr[M(x) = 0] \geq \frac{1}{2}$  and  $\Pr[M(x) = 1] = 0$ .
4. The expected running time of  $M$  on any input is polynomial.

Equivalently, ZPP is the class of languages that can be decided by Las Vegas algorithms: randomized algorithms that always give the correct answer when they terminate, and have polynomial expected running time.

**Statement.** Prove that  $\text{CL} \subseteq \text{ZPP}$ .

**4 Problem 4 (40 pts): Prove that  $\text{TC}^1 \subseteq \text{CL}$** 

In the class we sketch the high-level idea of the proof that  $\text{TC}^1 \subseteq \text{CL}$ . In this problem, you will complete the proof.

**Statement.** Prove that  $\text{TC}^1 \subseteq \text{CL}$ . You should give a complete register program for the  $\text{TC}^1$  circuit and prove it's correctness.

## Optional references and context

For accessible background on catalytic logspace and the toggling construction, see:

- Nathan Sheffield, *A quick-and-dirty intro to CL*.
- Ian Mertz, *Reusing Space: Techniques and Open Problems*.

### (Optional) Solution placeholders

You may use the following headings for your writeup; remove them if not needed.

## Solution to Problem 1

**Goal.** Prove that there exists a language  $L \in \text{SPACE}[n^2]$  such that for every constant  $c$ , for all sufficiently large  $n$ , no Boolean circuit of size at most  $cn$  computes  $L$  on  $\{0, 1\}^n$ ; equivalently,  $\text{SPACE}[n^2] \not\subseteq \text{SIZE}(O(n))$ .

**Proof.** We will proceed in steps.

**Step 1: Reduction to a hard function on  $2 \log n$  bits.**

Given input  $x \in \{0, 1\}^n$ . We are going to construct a function  $f_n : \{0, 1\}^{2 \log n} \rightarrow \{0, 1\}$  that is hard for circuits of size  $O(n^{1.1})$ . Our language  $L$  is going to be defined as  $L = \{x \in \{0, 1\}^n : f_n(x_{\leq 2 \log n}) = 1\}$ .

Note that if  $f_n$  is hard for circuits of size  $n^{1.1}$ , then  $L$  is hard for circuits of size  $O(n)$ .

**Step 2: Enumerating and testing functions.**

Fix  $m = 2 \log n$ . We will construct a function  $f : \{0, 1\}^m \rightarrow \{0, 1\}$  that is hard for circuits of size  $c \cdot m$  for a suitable constant  $c$ .

Our  $\text{SPACE}[n^2]$  algorithm does the following:

1. **Enumerate:** Iterate through all  $2^{2^m} = 2^{n^2}$  possible truth tables for functions  $f : \{0, 1\}^m \rightarrow \{0, 1\}$ , in lexicographic order. Each truth table can be represented as a string of length  $2^m = n^2$ , which fits in  $n^2$  space.
2. **Test hardness:** For each candidate function  $f$ , check whether  $f$  can be computed by some Boolean circuit of size at most  $n^{1.1}$ .

To perform this test in  $\text{SPACE}[n^2]$ :

- Enumerate all possible circuits  $C$  with at most  $n^{1.1}$  gates over the basis  $\{\wedge, \vee, \neg\}$  (or any complete basis).
  - For each circuit  $C$ , enumerate all  $2^m = n^2$  possible inputs  $x \in \{0, 1\}^m$  and check whether  $C(x) = f(x)$  for all  $x$ .
  - The number of circuits of size  $s$  is at most  $s^{O(s)}$ , so there are at most  $(n^{1.1})^{O(n^{1.1})} = 2^{O(n^{1.1} \log n)}$  circuits to check.
  - Each circuit evaluation takes polynomial space (in fact, logarithmic space suffices).
  - We can enumerate circuits using a counter of size  $O(n^{1.1} \log n)$  bits, which fits in  $n^2$  space for large  $n$ .
3. **Output:** Return the lexicographically first function  $f$  that is *not* computable by any circuit of size  $n^{1.1}$ .

**Step 3: Existence of a hard function.**

By a counting argument, such a hard function must exist. The number of Boolean functions on  $m = 2 \log n$  bits is  $2^{2^m} = 2^{n^2}$ . The number of circuits of size at most  $n^{1.1}$  is at most  $2^{O(n^{1.1} \log n)}$ . For sufficiently large  $n$ , we have

$$2^{O(n^{1.1} \log n)} \ll 2^{n^2},$$

so there must exist functions that cannot be computed by circuits of size  $n^{1.1}$ . Therefore, our algorithm will find such a function.

## Solution to Problem 2

**Claim.**  $\text{NEXP} \subseteq \text{coNEXP}/_{n+1}$ .

Let  $L \in \text{NEXP}$ . We let  $\alpha(n)$  be the number of yes-instances of  $L$  on inputs of length  $n$ . This can be specified in  $n + 1$  bits (since the range is from 0 to  $2^n$ ).

Let  $x \notin L$ . Given the correct advice  $\alpha(n)$ , we can verify that  $x \notin L$  as follows:

- We guess a list of  $\alpha(n)$  yes-instances of  $L$  on inputs of length  $n$ . For each of them we further guess their corresponding witnesses.
- We verify that all elements in this list are distinct, and that they are indeed yes-instances of  $L$  by verifying the witnesses. We reject immediately if we find a duplicate or a witness that is not correct.
- We then check that our  $x$  is not in the list. If it is, we reject and accept otherwise.

We first note that given the correct advice  $\alpha(n)$ , the above procedure would proceed to step 3 if and only if the guessed list  $L$  is the correct list of yes-instances of  $L$  on inputs of length  $n$ . Therefore, it would accept if and only if  $x \notin L$ . It also takes  $2^{O(n)}$  non-deterministic time.

Therefore, we have that  $L \in \text{coNEXP}/_{n+1}$ . Since  $L$  is arbitrary, we have that  $\text{NEXP} \subseteq \text{coNEXP}/_{n+1}$ .

□

## Solution to Problem 3

**Claim.**  $\text{CL} \subseteq \text{ZPP}$ .

Given an input  $x$ , we can think of the configuration of the machine as a pair  $(w, s)$ , where  $w \in \{0, 1\}^{O(\log n)}$  is the clean workspace/pointer locations/state of the machine, and  $s \in \{0, 1\}^{\text{poly}(n)}$  is the catalytic workspace/configuration of the machine.

Without loss of generality, we can assume the machine starts as  $w_0 = 0^{O(\log n)}$  but an arbitrary  $s_0 \in \{0, 1\}^{\text{poly}(n)}$ . By the definition of CL, the machine should halt with the same starting  $s_0$  in the catalytic workspace, but can have an arbitrary  $w$  in the clean workspace.

Let each pair  $(w, s)$  be a vertex, we can build a directed graph  $G$  where  $(w, s) \rightarrow (w', s')$  if and only if the machine transition from configuration  $(w, s)$  to configuration  $(w', s')$  in one step.

Then, for each starting configuration  $(w_0, s_0)$ , it corresponds to a path in  $G$  from  $(w_0, s_0)$  to some  $(w, s_0)$ .

Here we make the key claim that for two different starting configurations  $(w_0, s_0)$  and  $(w'_0, s'_0)$ , the corresponding paths in  $G$  are disjoint.

Indeed, if they were not disjoint, then they would intersect at some configuration  $(w, s)$ , and from thereafter the two paths would be the same. This contradicts the definition of CL since the machine should halt with the same starting  $s_0$  in the catalytic workspace.

Therefore, all these  $2^{|s|}$  paths must be disjoint. Since there are  $2^{|w|} \cdot 2^{|s|}$  many states in total, the expected length of a random path is at most  $2^{|w|} \cdot 2^{|s|}/2^{|s|} = 2^{|w|} \leq \text{poly}(n)$ .

Hence, our algorithm can work as follows:

1. Pick a random starting catalytic configuration  $s_0$ .
2. Simulate the machine starting from  $(w_0, s_0)$  until it halts and output the decision bit.

By the discussion above, the expected running time is at most  $\text{poly}(n)$ . Therefore, we have that  $\text{CL} \subseteq \text{ZPP}$ .

## Solution to Problem 4

**Claim.**  $\text{TC}^1 \subseteq \text{CL}$ .

**Proof.** Let  $L \in \text{TC}^1$ . Then for each input length  $n$ , there is a Boolean circuit  $C_n$  of depth  $O(\log n)$  and polynomial size over threshold gates deciding  $L$  on inputs of length  $n$ . Each threshold gate  $G$  in  $C_n$  has fan-in  $k = \text{poly}(n)$  and computes

$$G(p_1, \dots, p_k) = \mathbf{1} \left[ \sum_{i=1}^k p_i = T \right]$$

for some threshold  $T \in \{0, 1, \dots, k\}$ , where each  $p_i \in \{0, 1\}$  is the output of a gate at the previous layer. (Note that although  $\text{TC}^1$  is defined using majority gates, we can easily convert it to threshold gates by using the fact that

$$\text{MAJ}(p_1, \dots, p_k) = \sum_{T=\lceil k/2 \rceil}^k \mathbf{1} \left[ \sum_{i=1}^k p_i = T \right].$$

Our goal is to simulate the evaluation of  $C_n$  by a clean register program that:

- uses only  $O(\log n)$  *clean* bits (for the input, a program counter, and a recursion stack for the circuit of depth  $O(\log n)$ ), and
- uses at most  $\text{poly}(n)$  *dirty* bits, which are all restored to their initial values at the end of the computation.

This will show that  $L \in \text{CL}$ .

**Registers and invariant.** We work with registers  $R_1, R_2, \dots, R_M$  for some  $M = \text{poly}(n)$ , each holding an integer modulo a prime  $p$ . We will fix  $p$  so that it is larger than the number of gates in  $C_n$ . All arithmetic is performed modulo  $p$ . For each gate  $G$  of  $C_n$ , we will construct a clean register program  $P_G$  such that, given a designated output register  $R_D$ ,

- after running  $P_G(D)$ , all registers except  $R_D$  are restored to their original contents, and
- $R_D$  is increased by the Boolean output of the gate  $G$  (viewed as 0 or 1).

We will construct  $P_G(D)$  for each gate  $G$  in  $C_n$  in a bottom-up manner. First, for all input variables such programs can be constructed trivially. Now when we construct  $P_G(D)$  for a gate  $G$ , we will make use of the programs  $P_{G_1}(\cdot), \dots, P_{G_k}(\cdot)$  for the input gates  $G_1, \dots, G_k$ .

**Step 1: Summation of fan-in values.** Suppose  $G$  has fan-in  $k$ , with predecessor gates implemented by clean programs  $P_1, \dots, P_k$  (each  $P_i$  takes as argument the index of the register where it should add its output). Let  $p_i \in \{0, 1\}$  be the output of the  $i$ -th predecessor gate.

We define an auxiliary program  $P_{\text{sum}}$  that, given a register index  $D$ , temporarily computes  $\sum_{i=1}^k p_i$  and adds it into  $R_D$ , while restoring all other registers at the end:

$$P_{\text{sum}}(D) :$$

$$\forall i \in \{1, \dots, k\} : P_i(i)$$

$$R_D := R_D + \sum_{i=1}^k R_i$$

$$\forall i \in \{1, \dots, k\} : P_i^{-1}(i)$$

$$R_D := R_D - \sum_{i=1}^k R_i$$

Let  $C_i$  denote the initial value of register  $R_i$ , and  $C_D$  the initial value of  $R_D$ . By correctness and cleanliness of each  $P_i$ , after the first loop we have  $R_i = C_i + p_i$ . Thus after the assignment  $R_D := R_D + \sum_{i=1}^k R_i$ ,

$$R_D = C_D + \sum_{i=1}^k (C_i + p_i) = C_D + \sum_{i=1}^k C_i + \sum_{i=1}^k p_i.$$

After running  $P_i^{-1}$  for all  $i$ , each  $R_i$  is restored:  $R_i = C_i$ . Then

$$R_D := R_D - \sum_{i=1}^k R_i \Rightarrow R_D = C_D + \sum_{i=1}^k C_i + \sum_{i=1}^k p_i - \sum_{i=1}^k C_i = C_D + \sum_{i=1}^k p_i.$$

Thus all registers except  $R_D$  are restored, and  $R_D$  has been increased by  $\sum_{i=1}^k p_i$ .

**Step 2: Shifting by the threshold.** We next want to compute the difference  $T - \sum_{i=1}^k p_i$ . Let  $Q$  be an auxiliary register (dirty). We define a program  $P_{\text{aux}}$  that, given argument  $D$  and using  $Q$ , transforms  $R_D$  in such a way that the value  $T - \sum_i p_i$  can be accessed cleanly while keeping other registers restored at the end:

$P_{\text{aux}}(D)$  using register  $Q$ :

$$\begin{aligned} &P_{\text{sum}}(Q) \\ &R_D := R_D + T - R_Q \\ &P_{\text{sum}}^{-1}(Q) \\ &R_D := R_D + R_Q \end{aligned}$$

Tracing values, write  $C_Q$  for the initial content of  $R_Q$ . After  $P_{\text{sum}}(Q)$  we have

$$R_Q = C_Q + \sum_{i=1}^k p_i.$$

Thus after the update  $R_D := R_D + T - R_Q$  we get

$$R_D = C_D + T - (C_Q + \sum_{i=1}^k p_i) = C_D + T - C_Q - \sum_{i=1}^k p_i.$$

Now running  $P_{\text{sum}}^{-1}(Q)$  restores  $R_Q$  to  $C_Q$ , and leaves  $R_D$  unchanged. Finally, we do

$$R_D := R_D + R_Q$$

so that

$$R_D = C_D + T - C_Q - \sum_{i=1}^k p_i + C_Q = C_D + T - \sum_{i=1}^k p_i.$$

All other registers (including  $Q$ ) are restored. Thus  $P_{\text{aux}}$  cleanly adds the quantity  $T - \sum_i p_i$  into  $R_D$ .

### Step 3: Computing a high power.

Let  $\phi(p)$  be the Euler's totient function of  $p$ , note that  $\phi(p) = p - 1$  if  $p$  is prime. Importantly, we have that (this uses the assumption on  $p$  that  $p$  is larger than the number of gates in  $C_n$ )

$$\left( T - \sum_i p_i \right)^{\phi(p)} = 0 \iff T - \sum_i p_i = 0,$$

and

$$\left( T - \sum_i p_i \right)^{\phi(p)} = 1 \iff T - \sum_i p_i \neq 0.$$

We would like to compute a high power  $(T - \sum_i p_i)^{\phi(p)}$  for a suitable exponent  $\phi(p)$ , in a way that restores all registers except for a designated output register. Let us fix a prime  $p$  larger than an upper bound on  $T$  and  $k$ . We will compute  $(T - \sum_i p_i)^{\phi(p)}$  using a standard “toggling” trick.

Let  $Q$  be as in the previous step, and introduce auxiliary registers  $E, Q_1, \dots, Q_{\phi(p)}$ , all initially arbitrary (dirty). Define a program  $P_{\text{EXP}}$  that, given an input register  $D$ , uses these registers to add  $(T - \sum_i p_i)^{\phi(p)}$  to  $R_E$ , while restoring all registers except  $E$ :

$P_{\text{EXP}}(D)$  using registers  $Q, E, Q_1, \dots, Q_{\phi(p)}$  :

$$\begin{aligned} & P_{\text{aux}}^{-1}(Q) \\ & \forall i \in \{1, \dots, \phi(p)\} : R_{Q_i} := R_{Q_i} + R_Q^i \\ & P_{\text{aux}}(Q) \\ & \quad R_E := R_E + \sum_{i=0}^{\phi(p)} \binom{\phi(p)}{i} (-1)^i R_{Q_i} R_Q^{\phi(p)-i} \\ & P_{\text{aux}}^{-1}(Q) \\ & \forall i \in \{1, \dots, \phi(p)\} : R_{Q_i} := R_{Q_i} - R_Q^i \\ & P_{\text{aux}}(Q) \\ & \quad R_E := R_E - \sum_{i=0}^{\phi(p)} \binom{\phi(p)}{i} (-1)^i R_{Q_i} R_Q^{\phi(p)-i} \end{aligned}$$

To see correctness, let  $u := T - \sum_{i=1}^k p_i$ , and let  $C_Q, C_{Q_i}, C_E$  be the initial contents of the corresponding registers. One checks (using that  $P_{\text{aux}}$  and  $P_{\text{aux}}^{-1}$  add and subtract  $u$  cleanly) that the net effect on  $Q$  and each  $Q_i$  is zero, and that the increment to  $R_E$  is

$$R_E \leftarrow C_E + [(C_Q - u)^{\phi(p)} - C_Q^{\phi(p)}].$$

Using the binomial theorem and cancellation of the terms depending only on  $C_Q$ , this simplifies (independently of the initial dirty contents) to

$$R_E = C_E + u^{\phi(p)} = C_E + (T - \sum_{i=1}^k p_i)^{\phi(p)}.$$

All other registers are restored to their initial values.

**Step 4: Implementing a threshold gate.** We now use  $P_{\text{EXP}}$  to implement the gate  $G$ .

Note that we actually need to compute  $1 - (T - \sum_i p_i)^{\phi(p)}$  instead of  $(T - \sum_i p_i)^{\phi(p)}$  in order to get the correct output of the threshold gate. This is because the output of the threshold gate is 1 if and only if  $T - \sum_i p_i = 0$ .

To realize this as a clean register program, let  $Q$  be as before and let  $D$  be the designated output register for  $G$ . Define  $P_G(D)$  by:

$$P_G(D) :$$

$$\begin{aligned} & P_{\text{EXP}}(Q) \\ & R_D := R_D + 1 - R_Q \\ & P_{\text{EXP}}^{-1}(Q) \\ & R_D := R_D + R_Q \end{aligned}$$

Exactly as in the previous toggling arguments, one checks that all registers other than  $D$  are restored to their original values, while  $R_D$  is increased by  $1 - (T - \sum_i p_i)^{p-1} \in \{0, 1\}$ , i.e., by the output of gate  $G$ .

**Step 5: Simulating the whole circuit.** Starting from the input bits (which we load into designated registers), we construct a clean program for each gate in the circuit layer by layer. For each gate we use a fresh output register (dirty) to accumulate its output, treating all other registers as potentially dirty but required to be restored by the end of the gate's program. Because:

- each gate program uses only  $O(1)$  clean registers in addition to its input and output locations, and
- the circuit has depth  $O(\log n)$ ,

the total clean space needed to simulate the recursion/stack of the circuit evaluation is  $O(\log n)$ . At any point, we are using only a polynomial number of registers (dirty), each of size  $O(\log p) = O(\log n)$  bits, so the total dirty space is polynomial in  $n$ . By construction, at the end of the computation, all dirty registers are restored to their initial contents, and one designated output register holds the value of the circuit on the input  $x$ .

Thus we obtain a clean register program deciding  $L$  that uses  $O(\log n)$  clean space and polynomially many dirty bits, i.e.,  $L \in \text{CL}$ . Since  $L \in \text{TC}^1$  was arbitrary, we conclude that  $\text{TC}^1 \subseteq \text{CL}$ .  $\square$